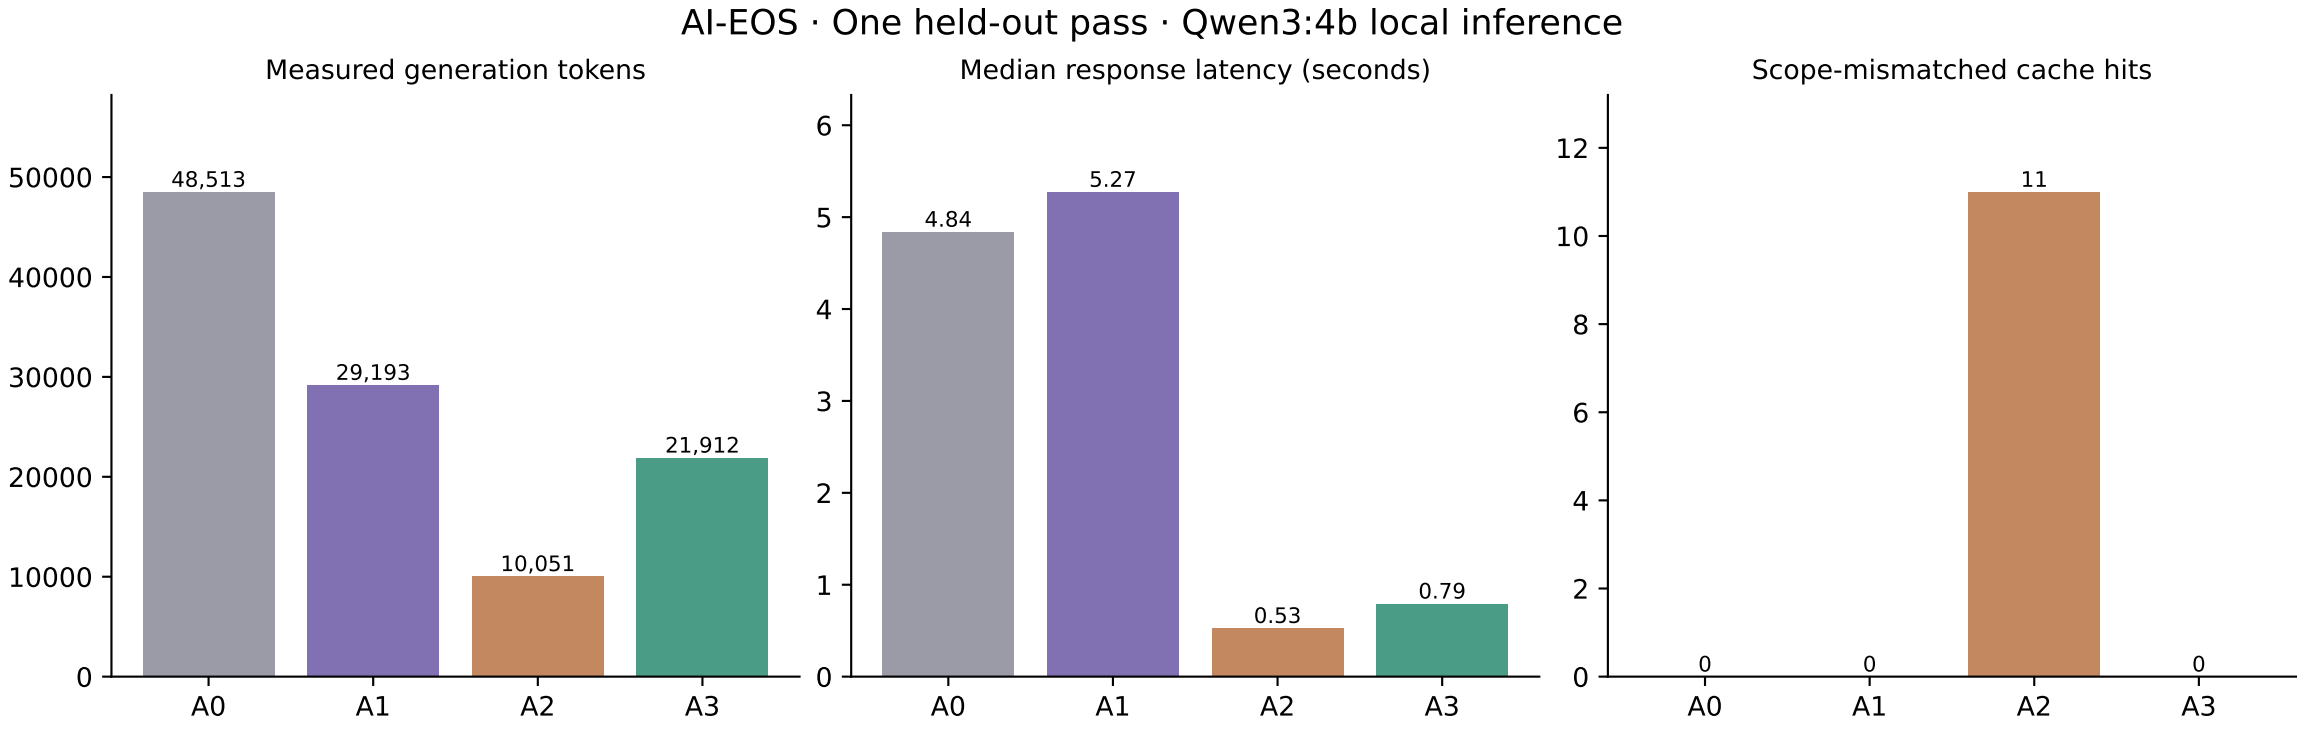
AI-EOS · One held-out pass · Qwen3:4b local inference
Measured generation tokens
48,513
29,193
10,051
21,912
0
10000
20000
30000
40000
50000
A0
A1
A2
A3
Median response latency (seconds)
4.84
5.27
0.53
0.79
0
1
2
3
4
5
6
A0
A1
A2
A3
Scope-mismatched cache hits
0
0
11
0
0
2
4
6
8
10
12
A0
A1
A2
A3